ECS 111 Homework Assignment #2

Kevin Oghalai

May 9, 2025

1 Introduction

A logistic regression model was used to predict the probability of a person having an income of over \$50,000. The dataset contained age, working class, education, marital status, occupation, and others. It additionally contained a label describing the salary of the person as either over or under \$50,000. The data was loaded in and then analyzed to determine possible issues that might arise from missing values and the various types of data. In this case, there were multiple missing values in addition to many categorical variables, which could not be directly used for analysis. To fix this, one-hot encoding was used. Although data could have been filled in with the mean or median of the column, this would not have working for categorical variables, and in the end, this did not end up being necessary.

2 Data Processing

This algorithm comes with options to regularize the data, in addition to loss functions, optimizers, and other hyperparameters. In this case, the dafault solver and loss function were attemped and then compared with another solver and loss function. The default loss and solver were L2 loss and the lbfgs solver. This was compared to the liblinear solver with L1 loss, then the liblinear solver with L2 loss. The results were then compared to the default values.

The data was split into a training and testing set. The model was trained on only the training set, and then the accuracy was determined on both sets later to investigate overfitting. As mentioned earlier one-hot encoding was used to deal with categorical variables.

3 Results

The logistic regression model had accuracy which was fairly dependent on the maximum number of iterations, so the results are plotted as a function of that. The results are shown in Figure 1.

Figure 1: Training accuracy as a function of the maximum number of iterations.

As shown in the figure above, the training accuracy ended up being higher when using the liblinear solver, although the loss function did not seem to matter. Unfortunately, the L1 loss function could not be used with the lbfgs solver, so the results are not shown. The training accuracy reached roughly 84.8% for the lbfgs solver, and 85.3% for the liblinear solver. Additionally, the lbfgs solver ended up being too slow to run for more than 10,000 iterations,even though it looks like it would go up more with more iterations. On the other hand, the liblinear solver was able to run much high numbers of iterations in a shorter time. The testing accuracy is shown below in Figure 2.

Figure 2: Training accuracy as a function of the maximum number of iterations.

Unlike above, the test accuracy actually went down with more iterations, which is likely due to overfitting. The test accuracy ended up going to 81% in all cases for a large number of iterations. The liblinear solver seemed to converge more quickly, but the lbfgs solver shows the same trend. In all cases with large numbers of iterations, the test accuracy was lower than the training accuracy, which is expected and shows that the data split was done correctly.

4 Discussion

While the logistic regression model was able to achieve a high accuracy, it is important to note that without actually evaluating what factors the model uses to determine the income, it is difficult to say why exactly the income is caused. This is a classic case of correlation vs. causation, where the model is able to predict the income, but it is not necessarily able to determine the cause of the income, leading to limited results being able to be drawn from the data. There is potential for misuse through ignorance or malice, as this is an issue which many people are likely to care about, but unable to actually analyze on their own. As an example, a reporter might say that getting married increases a person's salary, when in reality it might be that people with higher salaries are more likely to afford to get married. A slightly more humorous example would be someone buying a yacht because many billionaires own yachts. The yacht would be expensive and unlikely to increase the peron's salary, actually putting the person in debt, making them worse off.

In order to mitigate the issue, it would be useful to analyze exactly which factors cause

the model to predict higher incomes. In addition, experiments could be carried out tracking individuals throughout their lives in order to determine which factors are causing the income and which are simply correlated. This would be a difficult task, but would provide further insight and reduce bias.